

Suppose that you have trained a logistic regression classifier, and it outputs on a new example  $x$  a prediction  $h_{\theta}(x) = 0.7$ . This means (check all that apply):

1 point

- ☒ Our estimate for  $P(y = 1|x; \theta)$  is 0.7.
- ☒ Our estimate for  $P(y = 0|x; \theta)$  is 0.3.
- ☐ Our estimate for  $P(y = 0|x; \theta)$  is 0.7.
- ☐ Our estimate for  $P(y = 1|x; \theta)$  is 0.3.

Suppose you have the following training set, and fit a logistic regression classifier  $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$ .

1 point

$x_1$	$x_2$	$y$
1	0.5	0
1	1.5	0
2	1	1
3	1	0

Suppose that you have trained a logistic regression classifier, and it outputs on a new example  $x$  a prediction  $h_{\theta}(x) = 0.2$ . This means (check all that apply):

1 point

- ☐ Our estimate for  $P(y = 0|x; \theta)$  is 0.2.
- ☒ Our estimate for  $P(y = 0|x; \theta)$  is 0.8.
- ☒ Our estimate for  $P(y = 1|x; \theta)$  is 0.2.
- ☐ Our estimate for  $P(y = 1|x; \theta)$  is 0.8.

Suppose you have the following training set, and fit a logistic regression classifier  $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$ .

1 point


$x_1$	$x_2$	$y$
1	0.5	0
1	1.5	0
2	1	1
3	1	0

Suppose you have the following training set, and fit a logistic regression classifier

1 point

$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2).$

$x_1$	$x_2$	$y$
1	0.5	0
1	1.5	0
2	1	1
3	1	0



Which of the following are true? Check all that apply.

- ☒ Adding polynomial features (e.g., instead using  $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1 x_2 + \theta_5 x_2^2)$  ) could increase how well we can fit the training data.
- ☒ At the optimal value of  $\theta$  (e.g., found by fminunc), we will have  $J(\theta) \geq 0$ .
- ☐ Adding polynomial features (e.g., instead using  $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1 x_2 + \theta_5 x_2^2)$  ) would increase  $J(\theta)$  because we are now summing over more terms.

- ☐ If we train gradient descent for enough iterations, for some examples  $x^{(i)}$  in the training set it is possible to obtain  $h_{\theta}(x^{(i)}) > 1$ .

**For logistic regression, the gradient is given by  $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$ . Which of these is a correct gradient descent update for logistic regression with a learning rate of  $\alpha$ ? Check all that apply.**

1 point

- ☒  $\theta := \theta - \alpha \frac{1}{m} \sum_{i=1}^m \left( \frac{1}{1+e^{-\theta^T x^{(i)}}} - y^{(i)} \right) x^{(i)}.$
- ☐  $\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (\theta^T x - y^{(i)}) x_j^{(i)}$  (simultaneously update for all  $j$ ).
- ☒  $\theta := \theta - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}.$
- ☐  $\theta := \theta - \alpha \frac{1}{m} \sum_{i=1}^m (\theta^T x - y^{(i)}) x^{(i)}.$


**Which of the following statements are true? Check all that apply.**

1 point


- ☒ The cost function  $J(\theta)$  for logistic regression trained with  $m \geq 1$  examples is always greater than or equal to zero.
- ☐ Linear regression always works well for classification if you classify by using a threshold on the prediction made by linear regression.
- ☒ The sigmoid function  $g(z) = \frac{1}{1+e^{-z}}$  is never greater than one ( $> 1$ ).
- ☐ For logistic regression, sometimes gradient descent will converge to a local minimum (and fail to find the global minimum). This is the reason we prefer more advanced optimization algorithms such as fminunc (conjugate gradient/BFGS/L-BFGS/etc).

Suppose you train a logistic classifier  $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$ . Suppose  $\theta_0 = 6, \theta_1 = 0, \theta_2 = -1$ . Which of the following figures represents the decision boundary found by your classifier?

☐ Figure:



☐ Figure:



☒ Figure:

